

Emily Lyon

Deep Reinforcement Learning From Theory to Empirical Research

May 1, 2024

Springer Nature

Preface

Reinforcement learning is a branch of machine learning (ML), which is different from traditional machine learning such as supervised and unsupervised learning. It focused on learning from agent-environment interactions to achieve certain goal(s) optimally. The learning process is interactive and driven both internally and externally. Recent developments in other machine learning techniques, especially neural networks, have been driving forces which help advance this field significantly. The improvements in both the scope and the depth of problems being studied and solved in the area are encouraging. This book starts from introduction to math foundations, to recent improvements and advanced topics.

Why This Book

This book is written as a comprehensive introduction to the field of reinforcement learning, with the focus on recent improvements and new techniques in the area. It starts with the introduction to traditional reinforcement learning and its evolution, math foundations in the area, to the recent technique advances in the area that are being applied and developed, including deep reinforcement learning algorithms. Then, one chapter is dedicated to the realistic applications, followed by advanced topics in both academy and industry and attemptive solutions, which make RL-based models and systems empirical.

April, 2024

Emily Lyon, etc..

Contents

1	Introduction	1
1.1	Evolution of Reinforcement Learning	2
1.1.1	Basic Reinforcement Learning	2
1.1.2	Inverse Reinforcement Learning	2
1.1.3	Multi-agent Reinforcement Learning	2
1.1.4	Meta Reinforcement Learning	2
1.1.5	Hierarchical Reinforcement Learning	2
1.1.6	Multi-Task Reinforcement Learning	2
1.2	Applications of Reinforcement Learning	2
1.3	Deep Reinforcement Learning	2
1.4	Advanced Topics in Deep Reinforcement Learning	2
1.5	Summary	2
1.6	Bibliographic Notes	2
	References	2
2	Mathematics in Deep Reinforcement Learning	5
2.1	Reinforcement Learning Foundations	5
2.2	Deep Learning Foundations	5
2.3	Deep Reinforcement Learning Foundations	5
2.4	Performance Evaluations	5
2.4.1	Offline Performance Metrics	5
2.4.2	Online Performance Metrics	5
2.4.3	A/B Testing	5
2.4.4	Interleaving	5
	References	5
3	Deep Reinforcement Learning Models	7
3.1	General Functionalities of Deep Models	8
3.1.1	Policy Approximation Based	8
3.1.2	Value Function Approximation Based	8
3.2	Deep Embeddings	8

3.3	Transformers in Reinforcement Learning Models	8
3.4	CNN-based Reinforcement Learning Models	8
3.5	RNN-based Reinforcement Learning Models	8
3.6	Hybrid Deep Reinforcement Learning Models	8
3.7	Hybrid Reinforcement Learning Models	8
3.8	Advanced Network Components	8
3.8.1	Attention	8
	References	8
4	Emperical Deep Reinforcement Learning Systems	11
4.1	DRL In Robotics	11
4.2	DRL In Commerce	11
4.3	DRL In Medications	11
4.4	DRL In Education	11
4.5	DRL In Civilization	11
4.6	DRL In Simulations	11
	References	11
5	Performance Evaluations	13
5.0.1	Offline Performance Metrics	13
5.0.2	Online Performance Metrics	13
5.0.3	A/B Testing	13
5.0.4	Interleaving	13
	References	13
6	Advanced Topics	15
6.1	Memorization	15
6.2	Alignment of Target Network	15
6.3	Policy Drift and the Solutions	15
6.4	Sociality and Trust	15
6.5	Attack Resistance	15
6.6	Privacy Preservation	15
6.7	Personalization	15
	References	15
7	Empericcal Researches Ongoing	17
7.1	Adaptations of Neuroscience	17
7.2	Future of Deep Reinforcement Learning	17
	References	17
A	Appendix	19
A.1	Math Appendix	19
	Glossary	21
	Index	23

Acronyms

Use the template *acronym.tex* together with the document class SVMono (monograph-type books) or SVMult (edited books) to style your list(s) of abbreviations or symbols.

Lists of abbreviations, symbols and the like are easily formatted with the help of the Springer-enhanced `description` environment.

ABC	Spelled-out abbreviation and definition
BABI	Spelled-out abbreviation and definition
CABR	Spelled-out abbreviation and definition

Chapter 1

Introduction

Abstract Reinforcement Learning nowadays are widely used in various industries, including business, IT, Pharmacy and government, etc.. Along with the development and thrive of WWW and other important computer techniques, such as IOT and neural network, the breath and depth of the use of Reinforcement Learning have improved significantly. In this chapter, we highlight the evolution of Reinforcement Learning. Discuss the use of Reinforcement Learning in our daily life and beyond. Finally, we briefly discuss advanced topics in the area and leave the details to be elaborated in later chapters.

1.1 Evolution of Reinforcement Learning

1.1.1 Basic Reinforcement Learning

1.1.2 Inverse Reinforcement Learning

1.1.3 Multi-agent Reinforcement Learning

1.1.4 Meta Reinforcement Learning

1.1.5 Hierarchical Reinforcement Learning

1.1.6 Multi-Task Reinforcement Learning

1.2 Applications of Reinforcement Learning

1.3 Deep Reinforcement Learning

1.4 Advanced Topics in Deep Reinforcement Learning

1.5 Summary

1.6 Bibliographic Notes

References

In view of the parallel print and (chapter-wise) online publication of your book at www.springerlink.com it has been decided that – as a general rule – references should be sorted chapter-wise and placed at the end of the individual chapters. However, upon agreement with your contact at Springer you may list your references in a single separate chapter at the end of your book. Deactivate the class option `sectrefs` and the `thebibliography` environment will be put out as a chapter of its own.

References may be *cited* in the text either by number (preferred) or by author/year.¹ If the citation in the text is numbered, the reference list should be arranged in ascending order. If the citation in the text is author/year, the reference list should be

¹ Make sure that all references from the list are cited in the text. Those not cited should be moved to a separate *Further Reading* section or chapter.

sorted alphabetically and if there are several works by the same author, the following order should be used:

1. South J, Blass B (2001) The future of modern genomics. Blackwell, London

Chapter 2

Mathematics in Deep Reinforcement Learning

Abstract

2.1 Reinforcement Learning Foundations

2.2 Deep Learning Foundations

2.3 Deep Reinforcement Learning Foundations

2.4 Performance Evaluations

2.4.1 Offline Performance Metrics

2.4.2 Online Performance Metrics

2.4.3 A/B Testing

2.4.4 Interleaving

References

In view of the parallel print and (chapter-wise) online publication of your book at www.springerlink.com it has been decided that – as a general rule – references should be sorted chapter-wise and placed at the end of the individual chapters. However, upon agreement with your contact at Springer you may list your references in a single separate chapter at the end of your book. Deactivate the class option

sectrefs and the thebibliography environment will be put out as a chapter of its own.

References may be *cited* in the text either by number (preferred) or by author/year.¹ If the citation in the text is numbered, the reference list should be arranged in ascending order. If the citation in the text is author/year, the reference list should be *sorted* alphabetically and if there are several works by the same author, the following order should be used:

1. South J, Blass B (2001) The future of modern genomics. Blackwell, London

¹ Make sure that all references from the list are cited in the text. Those not cited should be moved to a separate *Further Reading* section or chapter.

Chapter 3

Deep Reinforcement Learning Models

3.1 General Functionalities of Deep Models

3.1.1 Policy Approximation Based

3.1.2 Value Function Approximation Based

3.2 Deep Embeddings

3.3 Transformers in Reinforcement Learning Models

3.4 CNN-based Reinforcement Learning Models

3.5 RNN-based Reinforcement Learning Models

3.6 Hybrid Deep Reinforcement Learning Models

3.7 Hybrid Reinforcement Learning Models

3.8 Advanced Network Components

3.8.1 Attention

References

In view of the parallel print and (chapter-wise) online publication of your book at www.springerlink.com it has been decided that – as a general rule – references should be sorted chapter-wise and placed at the end of the individual chapters.

However, upon agreement with your contact at Springer you may list your references in a single separate chapter at the end of your book. Deactivate the class option `sectrefs` and the `thebibliography` environment will be put out as a chapter of its own.

References may be *cited* in the text either by number (preferred) or by author/year.¹ If the citation in the text is numbered, the reference list should be arranged in ascending order. If the citation in the text is author/year, the reference list should be *sorted* alphabetically and if there are several works by the same author, the following order should be used:

1. South J, Blass B (2001) The future of modern genomics. Blackwell, London

¹ Make sure that all references from the list are cited in the text. Those not cited should be moved to a separate *Further Reading* section or chapter.

Chapter 4

Emperical Deep Reinforcement Learning Systems

4.1 DRL In Robotics

4.2 DRL In Commerce

4.3 DRL In Medications

4.4 DRL In Education

4.5 DRL In Civilization

4.6 DRL In Simulations

References

In view of the parallel print and (chapter-wise) online publication of your book at www.springerlink.com it has been decided that – as a genrerall rule – references should be sorted chapter-wise and placed at the end of the individual chapters. However, upon agreement with your contact at Springer you may list your references in a single seperate chapter at the end of your book. Deactivate the class option `sectrefs` and the `thebibliography` environment will be put out as a chapter of its own.

References may be *cited* in the text either by number (preferred) or by author/year.¹ If the citatiion in the text is numbered, the reference list should be arranged in ascending order. If the citation in the text is author/year, the reference list should be

¹ Make sure that all references from the list are cited in the text. Those not cited should be moved to a separate *Further Reading* section or chapter.

sorted alphabetically and if there are several works by the same author, the following order should be used:

1. South J, Blass B (2001) The future of modern genomics. Blackwell, London

Chapter 5

Performance Evaluations

5.0.1 Offline Performance Metrics

5.0.2 Online Performance Metrics

5.0.3 A/B Testing

5.0.4 Interleaving

References

In view of the parallel print and (chapter-wise) online publication of your book at www.springerlink.com it has been decided that – as a general rule – references should be sorted chapter-wise and placed at the end of the individual chapters. However, upon agreement with your contact at Springer you may list your references in a single separate chapter at the end of your book. Deactivate the class option `sectrefs` and the `thebibliography` environment will be put out as a chapter of its own.

References may be *cited* in the text either by number (preferred) or by author/year.¹ If the citation in the text is numbered, the reference list should be arranged in ascending order. If the citation in the text is author/year, the reference list should be *sorted* alphabetically and if there are several works by the same author, the following order should be used:

1. South J, Blass B (2001) The future of modern genomics. Blackwell, London

¹ Make sure that all references from the list are cited in the text. Those not cited should be moved to a separate *Further Reading* section or chapter.

Chapter 6

Advanced Topics

6.1 Memorization

6.2 Alignment of Target Network

6.3 Policy Drift and the Solutions

6.4 Sociality and Trust

6.5 Attack Resistance

6.6 Privacy Preservation

6.7 Personalization

References

In view of the parallel print and (chapter-wise) online publication of your book at www.springerlink.com it has been decided that – as a general rule – references should be sorted chapter-wise and placed at the end of the individual chapters. However, upon agreement with your contact at Springer you may list your references in a single separate chapter at the end of your book. Deactivate the class option `sectrefs` and the `thebibliography` environment will be put out as a chapter of its own.

References may be *cited* in the text either by number (preferred) or by author/year.¹

¹ Make sure that all references from the list are cited in the text. Those not cited should be moved to a separate *Further Reading* section or chapter.

If the citation in the text is numbered, the reference list should be arranged in ascending order. If the citation in the text is author/year, the reference list should be *sorted* alphabetically and if there are several works by the same author, the following order should be used:

1. South J, Blass B (2001) The future of modern genomics. Blackwell, London

Chapter 7

Empericcal Researches Ongoing

7.1 Adaptations of Neuroscience

7.2 Future of Deep Reinforcement Learning

References

In view of the parallel print and (chapter-wise) online publication of your book at www.springerlink.com it has been decided that – as a general rule – references should be sorted chapter-wise and placed at the end of the individual chapters. However, upon agreement with your contact at Springer you may list your references in a single separate chapter at the end of your book. Deactivate the class option `sectrefs` and the `thebibliography` environment will be put out as a chapter of its own.

References may be *cited* in the text either by number (preferred) or by author/year.¹ If the citation in the text is numbered, the reference list should be arranged in ascending order. If the citation in the text is author/year, the reference list should be *sorted* alphabetically and if there are several works by the same author, the following order should be used:

1. South J, Blass B (2001) The future of modern genomics. Blackwell, London

¹ Make sure that all references from the list are cited in the text. Those not cited should be moved to a separate *Further Reading* section or chapter.

Appendix A

Appendix

All's well that ends well

A.1 Math Appendix

Glossary

Use the template *glossary.tex* together with the Springer document class SVMono (monograph-type books) or SVMult (edited books) to style your glossary in the Springer layout.

glossary term Write here the description of the glossary term. Write here the description of the glossary term. Write here the description of the glossary term.

glossary term Write here the description of the glossary term. Write here the description of the glossary term. Write here the description of the glossary term.

glossary term Write here the description of the glossary term. Write here the description of the glossary term. Write here the description of the glossary term.

glossary term Write here the description of the glossary term. Write here the description of the glossary term. Write here the description of the glossary term.

glossary term Write here the description of the glossary term. Write here the description of the glossary term. Write here the description of the glossary term.

Index

acronyms, list of, ix
glossary, 21

symbols, list of, ix